

Efficient Algorithms

Ekkapot Charoenwanit

Software Systems Engineering

TGGS

KMUTNB

Lecture 5: Searching and Sorting

Binary Search Tree: Properties

Properties:

- (1) Each node x in a binary search tree (BST) has a key $x.key$.
- (2) Nodes x other than the root have parents $x.parent$.
- (3) Nodes x may have a left child $x.left$ or a right child $x.right$ or both. (**Branching Factor** = 2)

NB: These nodes are **pointers** unlike in a heap (from Lecture 3).

Binary Search Tree: Invariant

Invariant:

For any node x , for all nodes y in the **left** subtree of node x , $y.key \leq x.key$ and for all nodes z in the **right** subtree of node x , $z.key \geq x.key$.

Binary Search Tree: Search

The running time of the **search operation** is $O(h)$, where h is the height of the tree.

```
1: procedure BST-SEARCH(root, key)
2:   if root == NULL then
3:     return NULL
4:   else
5:     if key < root.key then
6:       return BST-SEARCH(root.left, key)
7:     else
8:       if key > root.key then
9:         return BST-SEARCH(root.right, key)
10:      else
11:        return root
12:
```

Binary Search Tree: Min and Max

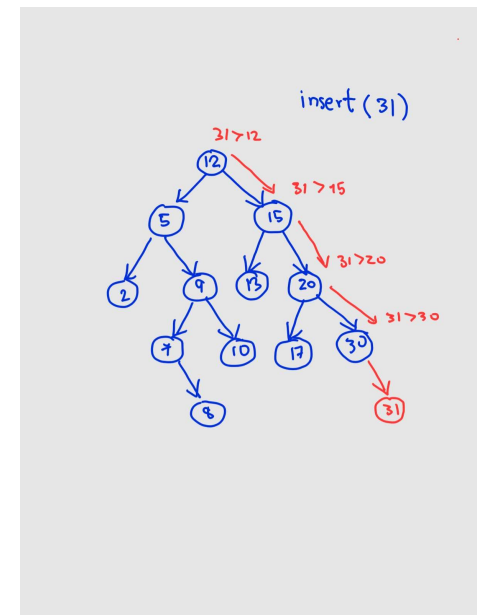
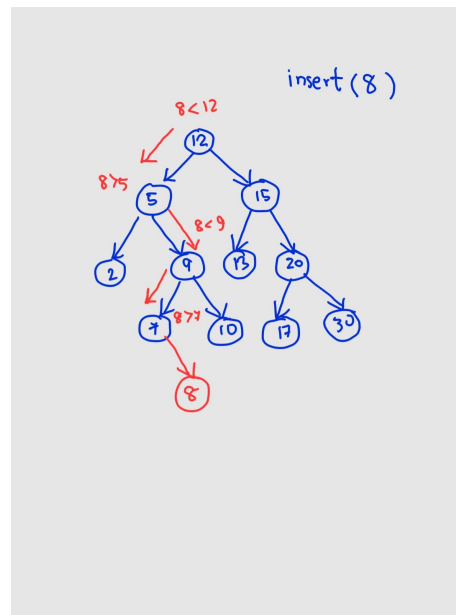
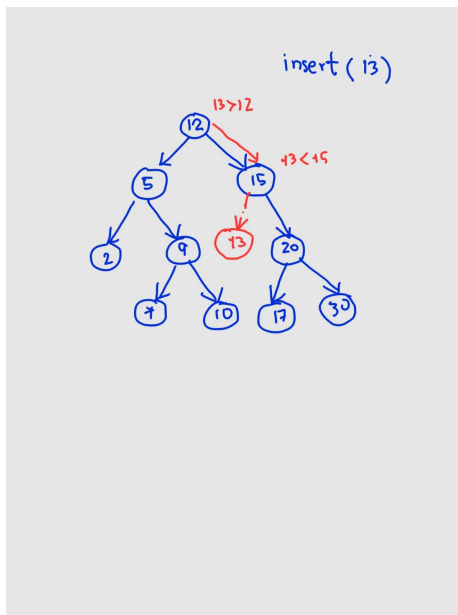
To find a **minimum key**,
keep going **left**.

To find a **maximum key**,
keep going **right**.

The running time of the *Min* and *Max* operations is the length of the downward path from root to the leftmost leaf and rightmost leaf, respectively.

Binary Search Tree: Insert

Follow the left and right pointers until the right position for the key being inserted is found.



Binary Search Tree: Delete

There are *three* cases.

Case I: The node x is a leaf.

- Follow the path until node x is reached.
- Remove x by modifying its parent to replace x with NULL as its child.

Binary Search Tree: Delete

There are **three** cases.

Case II: The node x has one child.

- Follow the path until x is reached.
- Remove x by elevating its child to take x 's position and modifying x 's parent to point to x 's child.

Binary Search Tree: Delete

There are **three** cases.

Case III: The node x has two children.

- Locate the leftmost leaf node y in the right subtree of x
- Replace x with y
- Remove the now duplicated leaf node y

The Sorting Problem

The sorting problem can be stated as follows:

Input: a sequence of n numbers $\langle a_1, a_2, \dots, a_n \rangle$,

Output: a permutation $\langle a_1', a_2', \dots, a_n' \rangle$ of the input sequence
 $a_1' \leq a_2' \leq \dots \leq a_n'$.

The numbers that we are sorting are also known as **keys**.

Insertion Sort

Idea:

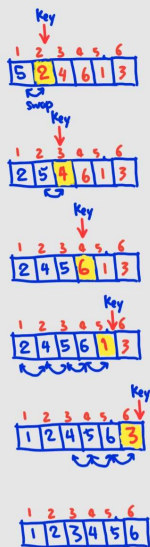
Imagine the way we sort a hand of playing cards.

- We start with an empty hand with the cards facing down on the table
- We pick up one card at a time from the table
- We insert the new card into the correct position in the left hand
 - To find the correct position, we compare the new card with the existing cards in the hand, from left to right.



*Illustration taken
from CLRS*

Insertion Sort



```
1: procedure INSERTION-SORT( $A$ )
2:   for  $j = 2 \rightarrow A.length$  do
3:      $key = A[j]$ 
4:      $i = j - 1$ 
5:     while  $i > 0 \wedge A[i] > key$  do
6:        $A[i + 1] = A[i]$ 
7:        $i = i - 1$ 
8:      $A[i + 1] = key$ 
```

Insertion Sort

The index j indicates the **current card** being inserted into the left hand.

Observation:

At the beginning of each iteration of the for loop,

- the subarray $A[1 \dots j - 1]$ constitutes the currently sorted hand
- the subarray $A[j + 1, \dots, n]$ corresponds to the pile of cards still on the table
- the elements $A[1 \dots j - 1]$ are the elements originally in positions 1 through $j - 1$, but now in sorted order.

Insertion Sort: Correctness Proof

Based on the *observation*,
we propose the following *loop invariant*:

At the start of each iteration of the for loop, the subarray $A[1 \dots j - 1]$ consists of the elements *originally* in $A[1 \dots j - 1]$ but in *sorted* order.

We will use this *loop invariant* to prove the *correctness* of insertion sort.

Loop Invariant

Initialization: The invariant is true prior to the first iteration of the loop.

Maintenance: If the invariant is true before an iteration of the loop, it remains true before the next iteration.

Termination: When the loop terminates, the invariant provides a useful property that can be used to show that the algorithm is correct.

Loop Invariant

When the **Initialization** & **Maintenance** properties are true, the loop invariant is true prior to every iteration.

This is the concept of **mathematical induction**:

- **Initialization** corresponds to the **base case**
- **Maintenance** corresponds to the **inductive step**
 - Assume true for an iteration i and show true for the next iteration i'
 - This is to show that the loop is inductive

The **Termination** property differs from how we use **mathematical induction** where we apply the inductive step **infinitely**.

At termination, we stop the induction and use the invariant to show that the algorithm is correct.

Insertion Sort: Loop Invariant

Initialization: Before the first iteration, $j = 2$.

The subarray $A[1 \dots 1] \equiv A[1]$ is ***trivially sorted*** and is also the original element in $A[1]$ so the invariant holds prior to the first iteration.

Maintenance: Informally, the body of the for loop works by moving $A[j - 1], A[j - 2], A[j - 3]$ and so on by one position to the right until it finds the proper position for $A[j]$, at which point $A[j]$ is inserted.

The subarray $A[1 \dots j]$ consists of the elements originally in $A[1 \dots j]$, but in sorted order.

When the loop counter j is incremented by one, the loop invariant still holds.

Insertion Sort: Termination

Termination: Finally, we check what happens when the loop terminates.

Because each iteration increments j by one, we have $j = n + 1$ on termination.

Substituting $j = n + 1$ into the wording of the loop invariant, we have that:

at the start of each iteration of the for loop, the subarray $A[1 \dots n]$, which is now the **entire array**, consists of the elements **originally** in $A[1 \dots n]$ but in **sorted** order. ■

Insertion Sort: Running Time Analysis

Barometer Instruction:

We count the total number of comparisons

- This corresponds to **Line 5**

```
1: procedure INSERTION-SORT( $A$ )
2:   for  $j = 2 \rightarrow A.length$  do
3:      $key = A[j]$ 
4:      $i = j - 1$ 
5:     while  $i > 0 \wedge A[i] > key$  do
6:        $A[i + 1] = A[i]$ 
7:        $i = i - 1$ 
8:      $A[i + 1] = key$ 
```

Insertion Sort: Worst-Case Analysis

At each iteration i of the **inner while loop**, the elements in $A[1 \dots i]$ are all moved one position to the right **in the worst case**.

- During each iteration, $i = j - 1$ **(Line 4)**
- There can be at most $\sum_{j=2}^n (j - 1)$ comparisons. **(Line 5)**

Summing up the number of comparisons:

$$\begin{aligned} T(n) &= \sum_{j=2}^n (j - 1) \\ &= \frac{n^2}{2} - \frac{n}{2} \blacksquare \end{aligned}$$

```
1: procedure INSERTION-SORT( $A$ )
2:   for  $j = 2 \rightarrow A.length$  do
3:      $key = A[j]$ 
4:      $i = j - 1$ 
5:     while  $i > 0 \wedge A[i] > key$  do
6:        $A[i + 1] = A[i]$ 
7:        $i = i - 1$ 
8:      $A[i + 1] = key$ 
```

Insertion Sort: Average-Case Analysis

For a given iteration i , there are i possible cases as follows:

Case	$A[i]$ is	#Comparisons
1	largest	1
2	second largest	2
3	third largest	3
...
$i - 1$	second smallest	$i - 1$
i	smallest	$i - 1$

Insertion Sort: Average-Case Analysis

Let N_i be a random variable representing the number of comparisons during iteration i .

$$E(N_i) = \sum_{k=1}^i n_k \Pr\{\text{Case } k \text{ happens}\}$$

$$E(N_i) = 1 \cdot \frac{1}{i} + 2 \cdot \frac{1}{i} + \cdots + (i-1) \cdot \frac{1}{i} + (i-1) \cdot \frac{1}{i}$$

$$E(N_i) = \frac{i-1}{2} + \frac{i-1}{i}$$

Insertion Sort: Average-Case Analysis

The *expected* total number N of comparisons, where $N = \sum_{i=2}^n N_i$, is

$$E(N) = \sum_{i=2}^n \left\{ \frac{i-1}{2} + \frac{i-1}{i} \right\} \leq \sum_{i=2}^n \frac{i-1}{2} + (n-1) \quad \left[\lim_{i \rightarrow \infty} \frac{i-1}{i} = 1 \right]$$

$$= \frac{n^2 + 3n - 4}{4}$$

Therefore, $E(N) = \Theta(n^2)$ ■

Insertion Sort: in-place algorithm

It is worth noting that insertion sort is an **in-place** algorithm in which $O(1)$ **auxiliary storage** is required to store **intermediate results** during the execution of the sorting algorithm.

Line 3 indicates that the variable *key* is the auxiliary space used by insertion sort.

```
1: procedure INSERTION-SORT(A)
2:   for  $j = 2 \rightarrow A.length$  do
3:      $key = A[j]$ 
4:      $i = j - 1$ 
5:     while  $i > 0 \wedge A[i] > key$  do
6:        $A[i + 1] = A[i]$ 
7:        $i = i - 1$ 
8:      $A[i + 1] = key$ 
```

Merge Sort

Idea: Merge sort is a **divide-and-conquer** algorithm

- Divide the array into **two subarrays** of (approximately the same size) if the array size is larger than one → **Divide**
- Sort each subarray (**recursively**) → **Conquer**
- Merge the **sorted** two subarrays → **Combine**

Merge Sort

Merge sort has several interesting properties as follows:

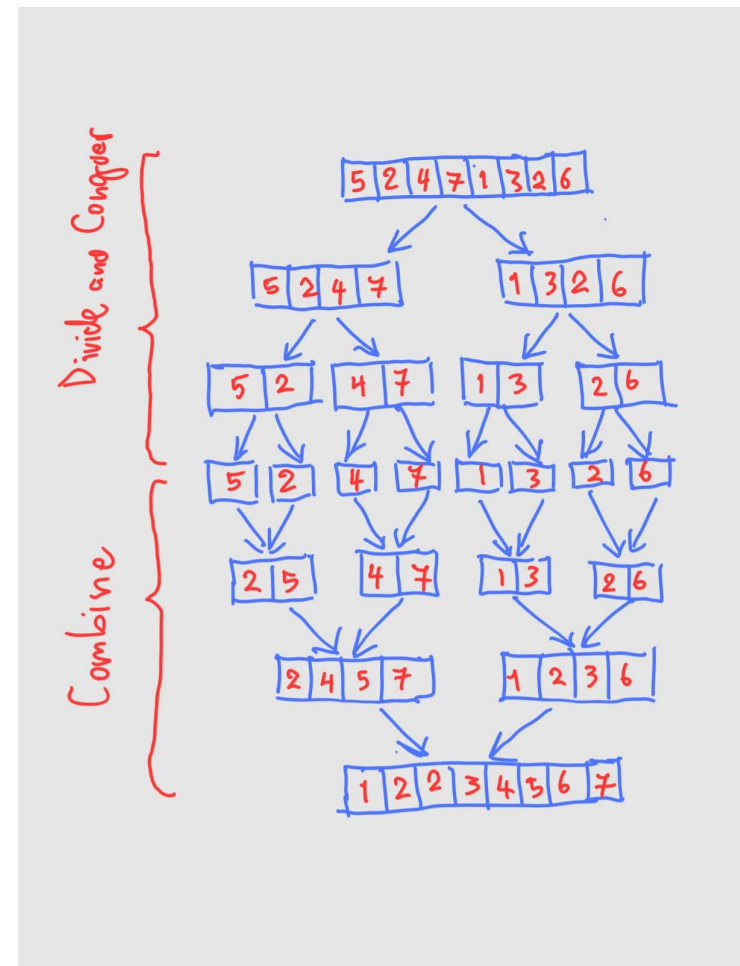
- The **divide** and **conquer** parts are very simple
- The **combine** part is (computationally) complex

```
1: procedure MERGE-SORT( $A, p, r$ )
2:   if  $p < r$  then
3:      $q = \lfloor \frac{p+r}{2} \rfloor$ 
4:     MERGE-SORT( $A, p, q$ )
5:     MERGE-SORT( $A, q + 1, r$ )
6:     MERGE( $A, p, q, r$ )
```

Merge Sort: Example

The **divide** and **conquer** parts are very simple

The **combine** part is (computationally) complex



Merge Sort: Merge

The *MERGE* routine is the **core part** of the **merge sort** algorithm as it is where most of the computational power is spent.

The *MERGE* assumes that the two subarrays $A[p \dots q]$ and $A[q + 1 \dots r]$ are **already sorted**.

Let n_1 and n_2 denote the number of elements in the **left-half** and the **right-half** subarrays of the array A , respectively. (**Lines 2-3**)

```
1: procedure MERGE( $A, p, q, r$ )
2:    $n_1 = q - p + 1$ 
3:    $n_2 = r - q$ 
4:   CREATE  $L[1 \dots n_1 + 1]$ 
5:   CREATE  $R[1 \dots n_2 + 1]$ 
6:   for  $i = 1 \rightarrow n_1$  do
7:      $L[i] = A[p + i - 1]$ 
8:   for  $i = 1 \rightarrow n_2$  do
9:      $R[i] = A[q + i]$ 
10:   $L[n_1 + 1] = \infty$ 
11:   $R[n_2 + 1] = \infty$ 
12:   $i = 1$ 
13:   $j = 1$ 
14:  for  $k = p \rightarrow r$  do
15:    if  $L[i] \leq R[j]$  then
16:       $A[k] = L[i]$ 
17:       $i = i + 1$ 
18:    else
19:       $A[k] = R[j]$ 
20:       $j = j + 1$ 
```

Merge Sort: Merge

Auxiliary arrays L and R are created to hold the left-half and the right-half subarrays, respectively. (**Lines 4-9**)

The so called **two-finger** algorithm is used to compare and merge the two subarrays. (**Lines 14-20**)

- The **left finger** points to the smallest element in the **left subarray** L that has not been copied back to A
- The **right finger** points to the smallest element in the **right subarray** R that has not been copied back to A

```
1: procedure MERGE( $A, p, q, r$ )
2:    $n_1 = q - p + 1$ 
3:    $n_2 = r - q$ 
4:   CREATE  $L[1 \dots n_1 + 1]$ 
5:   CREATE  $R[1 \dots n_2 + 1]$ 
6:   for  $i = 1 \rightarrow n_1$  do
7:      $L[i] = A[p + i - 1]$ 
8:   for  $i = 1 \rightarrow n_2$  do
9:      $R[i] = A[q + i - 1]$ 
10:   $L[n_1 + 1] = \infty$ 
11:   $R[n_2 + 1] = \infty$ 
12:   $i = 1$ 
13:   $j = 1$ 
14:  for  $k = p \rightarrow r$  do
15:    if  $L[i] \leq R[j]$  then
16:       $A[k] = L[i]$ 
17:       $i = i + 1$ 
18:    else
19:       $A[k] = R[j]$ 
20:       $j = j + 1$ 
```

Merge Sort: Proof of Correctness

Since *MERGE* is the core part, we will focus on the correctness of the *MERGE* routine by considering its **loop invariant**.

Once we show the correctness of *MERGE*, it is easy to see the correctness of *MergeSort*.

Loop Invariant:

- (1) At the start of each iteration of the for loop of **Lines 14-20**, the subarray $A[p \dots k - 1]$ contains the $k - p$ smallest elements of $L[1 \dots n_1 + 1]$ and $R[1 \dots n_2 + 1]$ in sorted order.
- (2) $L[i]$ and $R[j]$ are the smallest elements of the two arrays that have not been copied back to A .

Merge Sort: Proof of Correctness

Initialization:

Prior to the first iteration, $k = p$.

Therefore, the subarray $A[p \dots p - 1]$ is **empty** and contains $k - p = k - k = 0$ smallest elements.

Moreover, $i = j = 1$. $L[1]$ and $R[1]$ contain the smallest elements of the two subarrays that have not been copied back into A .

Merge Sort: Proof of Correctness

Maintenance:

Consider the case $L[i] \leq R[j]$.

By the loop invariant and the fact that $L[i] \leq R[j]$, $L[i]$ is the smallest element not copied back to A .

By the loop invariant, $A[p \dots k-1]$ contains the smallest $k-p$ elements.

By inspection, $A[k] = L[i]$ (**Line 16**).

(1) Therefore, $A[p \dots k]$ **now** contains the smallest $k-p+1$ elements.

(2) Moreover, $L[i+1]$ ($L[i]$ has just been copied back to A by **Line 16**) and $R[j]$ contain the smallest elements of the two subarrays that have not been copied back to A .

Merge Sort: Proof of Correctness

Maintenance: (Continued)

After **Line 17**, i is incremented, j remains the same and k is incremented:

$$i' = i + 1, j' = j \text{ and } k' = k + 1$$

(1) Therefore, $A[p \dots k' - 1]$ contains the smallest $k' - p$ elements.

(2) $L[i']$ and $R[j']$ now contains the smallest elements of the two arrays that have not been copied back to A .

This reestablishes the loop invariant at the start of the next iteration.

***The maintenance property for the other case $L[i] > R[j]$ can be shown in a similar way.

Merge Sort: Proof of Correctness

Termination:

When the for loop terminates, $k = r + 1$.

By the definition of the loop invariant,

$A[p \dots k - 1] = A[p \dots (r + 1) - 1] = A[p \dots r]$ contains the smallest $k - p = (r + 1) - p = r - p + 1$ elements in sorted order.

Therefore, we have just shown that *MERGE* works correctly by producing a sorted array $A[p \dots r]$ ■

Merge Sort: Time Complexity

The running time $T(n)$ for the base case ($n = 1$) is $O(1)$.

The running time $T(n)$ for the recursive case ($n > 1$) is split into the following components:

the **divide** part costs $O(1)$

the **conquer** part costs $2T\left(\frac{n}{2}\right)$

the **combine** part costs $O(n)$

Therefore, $T(n) = O(1) + 2T\left(\frac{n}{2}\right) + O(n)$

Allowing the asymptotically smaller term $O(1)$ to be absorbed into the asymptotically larger term $O(n)$, we have

$$T(n) = 2T\left(\frac{n}{2}\right) + O(n) \text{ for } n > 1 \blacksquare$$

Merge Sort: Time Complexity

Solving the recurrence relation, we get

$$T(n) = \Theta(n \log n)$$

NB: you solved this using the *recursion tree method* in PS2.4 as homework.

Merge Sort: out-of-place

Merge sort needs *auxiliary space* $\Theta(n)$ to store *intermediate results* during the execution of the sorting algorithm.

Therefore, it is *not* an in-place algorithm.

Summary

We have covered the following topics:

- Binary Search Tree and its basic operations
- Insertion Sort
- Merge Sort
- Correctness Proof using Loop Invariants

In the next lecture, we will cover *divide and conquer*.