

Efficient Algorithms

Ekkapot Charoenwanit

Software Systems Engineering

TGGS

KMUTNB

Lecture 11: Graph Algorithms (Part III)

All-Pair Shortest Paths (APSP)

Shortest Path Problem

In a shortest path problem, we are given a **weighted, directed** graph $G = (V, E, w)$ with a weight function $w: E \rightarrow \mathbb{R}$ that maps edges to **real-valued** weights.

The weight $w(p)$ of a path $p = \langle v_0, v_1, v_2, \dots, v_k \rangle$ is the sum of the weights of its constituent edges:

$$w(p) = \sum_{i=1}^k w(v_{i-1}, v_i)$$

Shortest Path Problem

We define the **shortest path weight** from u to v by

$$\delta(u, v) = \begin{cases} \min\{w(p) : u \rightsquigarrow v\} \\ \infty \end{cases} \quad (\text{Eq.1})$$

A **shortest path** from u to v is then defined as **any path** p with weight $w(p) = \delta(u, v)$.

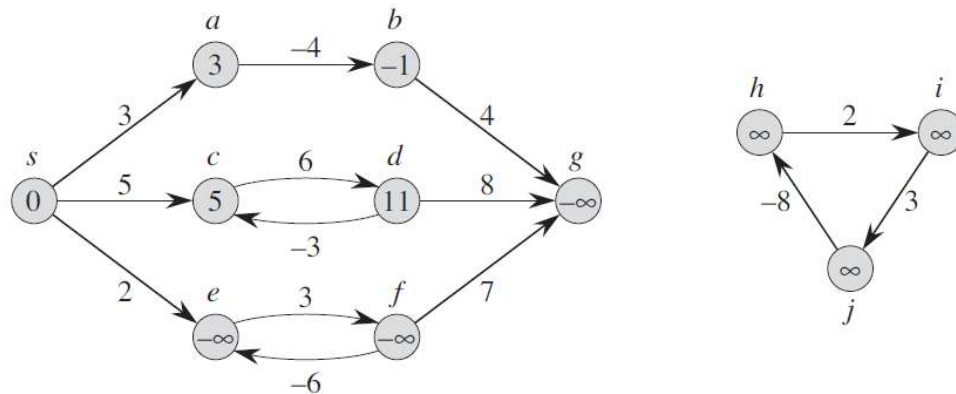
If there is no path from u to v , $\delta(u, v) = \infty$.

Even though there is a path u to v , a shortest path may not exist in the presence of **a negative-weight cycle** reachable from u .

Negative-Weight Cycle

Even though there is a path u to v , a shortest path may not exist in the presence of at least one **negative-weight cycle** reachable from u . Thus, $\delta(u, v) = -\infty$.

In the example below, a shortest path from s to f is **undefined** because we can always find a path with a smaller weight by traversing the negative-weight cycle $\langle e, f, e \rangle$ as many times as we want before reaching f .



All-Pair Shortest Paths: Ad-hoc Solutions

If $G = (V, E)$ contains **no negative-weight edges**, we can run **Dijkstra's Algorithm** for each $v \in V$ as **source vertex**.

- The algorithm runs in $O((V + E) \log V) \cdot |V| = O((V^2 + VE) \log V)$ time.
- If G is **dense**, the running time is $O(V^3 \log V)$.

If $G = (V, E)$ contains **negative-weight edges**, we can run **Bellman-Ford** for each $v \in V$ as **source vertex**.

- The algorithm runs in $O(VE) \cdot |V| = O(V^2 E)$ time.
- If G is **dense**, the running time is $O(V^4)$.

We will see that we can achieve a **better time complexity** than that of these two ad-hoc solutions.

DP Solution

Assuming that a weighted, directed graph $G = (V, E, w)$ is represented by an **adjacency matrix** $W = (w_{ij})$, consider a **shortest path** p from a vertex i to a vertex j .

Suppose that p contains **at most** m edges, and if there are **no negative-weight cycles** on p , m is **finite**.

If $i = j$, $w(p) = 0$ and p contains **no edges**.

If $i \neq j$, we can break p into $i \rightsquigarrow k \rightarrow j$, where $p' = i \rightsquigarrow k$, and p' contains **at most** $m - 1$ edges.

- Moreover, p' is a shortest path from i to k by the **Optimal Substructure Lemma** (See **Lecture 10**).

DP Solution: Recurrence Formulation

Let $l_{ij}^{(m)}$ be the **minimum weight** of any path from vertex i to vertex j that contains **at most** m edges.

Base Case: when $m = 0$, we have $l_{ij}^{(0)} = \begin{cases} 0 & \text{if } i = j \\ \infty & \text{if } i \neq j \end{cases}$

That is, there is only a shortest path if and only if $i = j$.

Recursive Case: we **exhaustively** search for the minimum among all possible predecessors k of j .

$$\begin{aligned} l_{ij}^{(m)} &= \min(l_{ij}^{(m-1)}, \min_{1 \leq k \leq n} (l_{ik}^{(m-1)} + w(k, j))) \\ &= \min_{1 \leq k \leq n} (l_{ik}^{(m-1)} + w(k, j)) \quad [w_{jj} = 0 \text{ for all } j] \end{aligned} \quad (\text{Eq.2})$$

DP Solution

Observation: If G contains **no negative-weight cycles**, then for every pair of vertices i and j for which $\delta(i, j) < \infty$, there is a **shortest path** from i to j that is **simple** with **at most** $n - 1$ edges.

This **observation** implies the following equalities:

$$\delta(i, j) = l_{ij}^{(n-1)} = l_{ij}^{(n)} = l_{ij}^{(n+1)} = \dots$$

DP Solution: Bottom-Up Solution

Taking as input the adjacency matrix $W = (w_{ij})$, we compute a series of matrices $L^{(1)}, L^{(2)}, L^{(3)}, \dots, L^{(n-1)}$, where $L^{(m)} = (l_{ij}^{(m)})$ for $i = 1, 2, 3 \dots, n - 1$.

The **final matrix** $L^{(n-1)}$ contains the **actual shortest path weights** for every pair of vertices.

In essence, the algorithm on the right, given $L^{(m-1)}$ and W as input, produces $L^{(m)}$ as output, effectively extending the shortest paths computed so far by **one more edge**:

$L^{(0)}$ and W and produces $L^{(1)}$,

$L^{(1)}$ and W and produces $L^{(2)}$,

....

$L^{(n-2)}$ and W and produces $L^{(n-1)}$

EXTEND-SHORTEST-PATHS(L, W)

```
1   $n = L.rows$ 
2  let  $L' = (l'_{ij})$  be a new  $n \times n$  matrix
3  for  $i = 1$  to  $n$ 
4      for  $j = 1$  to  $n$ 
5           $l'_{ij} = \infty$ 
6      for  $k = 1$  to  $n$ 
7           $l'_{ij} = \min(l'_{ij}, l_{ik} + w_{kj})$ 
8  return  $L'$ 
```

DP Solution: Bottom-Up Solution

Since the algorithm consists of **triplely nested** for loops, each of which executes for exactly n times, the running time is $\Theta(n^3)$.

EXTEND-SHORTEST-PATHS(L, W)

```
1   $n = L.rows$ 
2  let  $L' = (l'_{ij})$  be a new  $n \times n$  matrix
3  for  $i = 1$  to  $n$ 
4      for  $j = 1$  to  $n$ 
5           $l'_{ij} = \infty$ 
6          for  $k = 1$  to  $n$ 
7               $l'_{ij} = \min(l'_{ij}, l_{ik} + w_{kj})$ 
8  return  $L'$ 
```

DP Solution: Matrix Multiplication View

We will show that the DP solution we have just arrived at has a close relation to **matrix multiplication**.

Recall that, to compute $C = A \cdot B$ of two $n \times n$ matrices, we compute the following:

$$c_{ij} = \sum_{k=1}^n a_{ik} \cdot b_{kj} \text{ for } i, j = 1, 2, \dots, n \quad (\text{Eq.3})$$

Observe that if we make the following **symbolic substitutions in Eq.2**, we obtain **Eq.3**.

$$\begin{array}{ll} l^{(m-1)} & \rightarrow a \\ w & \rightarrow b \\ l^{(m)} & \rightarrow c \\ \min & \rightarrow + \\ + & \rightarrow \cdot \end{array}$$

DP Solution: Matrix Multiplication View

Thus, if we apply these changes to *Extend – Shortest – Paths* and replace ∞ with 0, we obtain the following matrix multiplication algorithm we are familiar with. (See **Lecture 6**)

```
SQUARE-MATRIX-MULTIPLY( $A, B$ )
1   $n = A.rows$ 
2  let  $C$  be a new  $n \times n$  matrix
3  for  $i = 1$  to  $n$ 
4      for  $j = 1$  to  $n$ 
5           $c_{ij} = 0$ 
6          for  $k = 1$  to  $n$ 
7               $c_{ij} = c_{ij} + a_{ik} \cdot b_{kj}$ 
8  return  $C$ 
```

DP Solution: Matrix Multiplication View

Let's return to the **APSP** Problem.

Recall that we compute the shortest-path weights by extending shortest paths **edge by edge**.

Letting $A \cdot B$ denote the matrix "**product**" returned by *Extend – Shortest – Paths*(A, B), we can compute the sequence of $n - 1$ matrices as follows:

$$L^{(1)} = L^{(0)} \cdot W = W$$

$$L^{(2)} = L^{(1)} \cdot W = W^2$$

$$L^{(3)} = L^{(2)} \cdot W = W^3$$

$$L^{(n-1)} = L^{(n-2)} \cdot W = W^{n-1}$$

Recall that the **final matrix** $L^{(n-1)}$ contains the **actual shortest path weights** for every pair of vertices.

DP Solution: Matrix Multiplication View

We arrive at the following **SLOW** algorithm to compute all-pair shortest paths.

The algorithm starts by initializing $L^{(1)} = W$ and extends the shortest paths computed so far **by one edge at a time**.

- $L^{(n-1)}$ returned by the algorithm contains the actual shortest paths.

The algorithm runs in $\Theta(n^4)$ time because there are $n - 2 = \Theta(n)$ iterations, each of which runs in $\Theta(n^3)$ time.

SLOW-ALL-PAIRS-SHORTEST-PATHS(W)

```
1   $n = W.rows$ 
2   $L^{(1)} = W$ 
3  for  $m = 2$  to  $n - 1$ 
4      let  $L^{(m)}$  be a new  $n \times n$  matrix
5       $L^{(m)} = \text{EXTEND-SHORTEST-PATHS}(L^{(m-1)}, W)$ 
6  return  $L^{(n-1)}$ 
```

DP Solution: Matrix Multiplication View

We can apply *the repeated squaring technique* to reduce the number of matrix multiplications to $\lceil \log_2(n-1) \rceil$ as follows:

$$L^{(1)} = W$$

$$L^{(2)} = W^2 = W \cdot W$$

$$L^{(4)} = W^4 = W^2 \cdot W^2$$

$$L^{(8)} = W^8 = W^4 \cdot W^4$$

$$L^{(2^{\lceil \log_2(n-1) \rceil})} = W^{2^{\lceil \log_2(n-1) \rceil}} = W^{2^{\lceil \log_2(n-1) \rceil - 1}} \cdot W^{2^{\lceil \log_2(n-1) \rceil - 1}}$$

Since $2^{\lceil \log_2(n-1) \rceil} \geq n-1$, it is always guaranteed that

$$L^{(2^{\lceil \log_2(n-1) \rceil})} = L^{(n-1)}$$

DP Solution: Matrix Multiplication View

The following **faster** algorithm can achieve a better time complexity of $\Theta(n^3 \log n)$ using the repeated squaring technique:

- There are $\lceil \log_2(n - 1) \rceil$ matrix multiplications, each of which takes $\Theta(n^3)$ time.

FASTER-ALL-PAIRS-SHORTEST-PATHS(W)

```
1   $n = W.rows$ 
2   $L^{(1)} = W$ 
3   $m = 1$ 
4  while  $m < n - 1$ 
5      let  $L^{(2m)}$  be a new  $n \times n$  matrix
6       $L^{(2m)} = \text{EXTEND-SHORTEST-PATHS}(L^{(m)}, L^{(m)})$ 
7       $m = 2m$ 
8  return  $L^{(m)}$ 
```

The Floyd-Warshall Algorithm

In the **Floyd-Warshall** algorithm, a different way of characterizing the structure of shortest paths is used.

The **Floyd-Warshall** algorithm considers the **intermediate vertices** of a shortest path, where an intermediate vertex of a **simple path** $p = \langle v_1, v_2, \dots, v_l \rangle$ is any vertex of p other than v_1 or v_l .

- That is, it can be any vertex in the set $\{v_2, v_3, \dots, v_{l-1}\}$.

The Floyd-Warshall Algorithm

For any pair of vertices $i, j \in V$, consider all paths from i to j with **all intermediate vertices** drawn from the set $\{1, 2, \dots, k\}$ and let p be a **minimum weight-path** among all of them.

Case I: If k is **not** an intermediate vertex of path p , then all intermediate vertices of p are in the set $\{1, 2, \dots, k-1\}$.

- A shortest path from i to j with all intermediate vertices in the set $\{1, 2, \dots, k-1\}$ is also a shortest path from i to j with all intermediate vertices in the set $\{1, 2, \dots, k\}$.

Case II: If k is an intermediate vertex of path p , then we decompose p into $i \rightsquigarrow k \rightsquigarrow j$: $p_1 = i \rightsquigarrow k$ and $p_2 = k \rightsquigarrow j$.

- By the **Optimal Substructure Property**, p_1 is a shortest path from i to k with all the intermediate vertices in the set $\{1, 2, \dots, k\}$.
- In fact, we can make **a stronger statement**, because k is not an intermediate vertex of p_1 , all intermediate vertices of p_1 are in the set $\{1, 2, \dots, k-1\}$.
- Therefore, p_1 is a shortest path from i to k with all intermediate vertices in the set $\{1, 2, \dots, k-1\}$.
- Similarly, p_2 is a shortest path from k to j with all intermediate vertices in the set $\{1, 2, \dots, k-1\}$.

The Floyd-Warshall Algorithm

Let $d_{ij}^{(k)}$ be the weight of a shortest path from i to j for which all the intermediate vertices are in the set $\{1, 2, \dots, k\}$.

$d_{ij}^{(0)}$ is a shortest path from i to j with no intermediate vertices at all so such a path contains **at most one edge**.

$$d_{ij}^{(k)} = \begin{cases} w_{ij} & \text{if } k = 0 \\ \min \left(d_{ij}^{(k-1)}, d_{ik}^{(k-1)} + d_{kj}^{(k-1)} \right) & \text{if } k \geq 1 \end{cases} \quad (\text{Eq.4})$$

Because, for any path, all intermediate vertices are in the set $\{1, 2, \dots, n\}$, the matrix $D^{(n)} = (d_{ij}^{(n)})$ is the **final solution** where $\delta(i, j) = d_{ij}^{(n)}$ for all $i, j \in V$.

The Floyd-Warshall Algorithm

The Floyd-Warshall Algorithm consists of **three nested for loops**, each of which executes **exactly** n iterations.

Because each execution of **line 7** takes $\Theta(1)$ time, the total running time is $n^3 \cdot \Theta(1) = \Theta(n^3)$.

```
FLOYD-WARSHALL( $W$ )
1   $n = W.rows$ 
2   $D^{(0)} = W$ 
3  for  $k = 1$  to  $n$ 
4      let  $D^{(k)} = (d_{ij}^{(k)})$  be a new  $n \times n$  matrix
5      for  $i = 1$  to  $n$ 
6          for  $j = 1$  to  $n$ 
7               $d_{ij}^{(k)} = \min(d_{ij}^{(k-1)}, d_{ik}^{(k-1)} + d_{kj}^{(k-1)})$ 
8  return  $D^{(n)}$ 
```

Transitive Closure of a Directed Graph

Given a **weighted, directed** graph $G = (V, E)$ with vertex set $V = \{1, 2, \dots, n\}$, we can determine whether G contains a path from i to j for all pair of vertices $i, j \in V$.

We define the **transitive closure** of G as $G^* = (V, E^*)$, where $E^* = \{(i, j): \text{there is a path from } i \text{ to } j \text{ in } G\}$.

The most obvious way to compute the **transitive closure** of a graph is to assign a weight of 1 to each edge of E and run **the Floyd-Warshall algorithm** on the graph.

- If there is a path from i to j , $d_{ij} < n$.
- Otherwise, $d_{ij} = \infty$.
- The algorithm takes $\Theta(n^3)$.

Transitive Closure of a Directed Graph

Another method for computing the **transitive closure** is to rely on the **matrix multiplication view**:

$$\begin{array}{rcl} \min & \rightarrow & \vee \\ + & \rightarrow & \wedge \end{array}$$

Let $t_{ij}^{(k)}$ to be

- 1 if there is a path from i to j with all intermediate vertices in the set $\{1, 2, \dots, k\}$
- 0 otherwise

Thus, we can formulate a recurrence solution as follows:

Base Case:

$$t_{ij}^{(0)} = \begin{cases} 0 & \text{if } i \neq j \wedge (i, j) \notin E \\ 1 & \text{if } i = j \vee (i, j) \in E \end{cases} \quad (\text{Eq.5})$$

Recursive Case:

$$t_{ij}^{(k)} = t_{ij}^{(k-1)} \vee (t_{ik}^{(k-1)} \wedge t_{kj}^{(k-1)}) \quad \text{for } k \geq 1 \quad (\text{Eq.6})$$

Transitive Closure of a Directed Graph

As in the Floyd-Warshall algorithm, we compute the sequence of matrices $T^{(k)} = (t_{ij}^{(k)})$ for $k = 1, 2, \dots, n - 1$.

The algorithm takes $\Theta(n^3)$ time.

- There are **three nested for loops**, each of which executes **exactly** n iterations.
- Therefore, there are n^3 iterations in total.
- Each iteration takes $\Theta(1)$ time (**line 12**)

```
TRANSITIVE-CLOSURE( $G$ )
1   $n = |G.V|$ 
2  let  $T^{(0)} = (t_{ij}^{(0)})$  be a new  $n \times n$  matrix
3  for  $i = 1$  to  $n$ 
4      for  $j = 1$  to  $n$ 
5          if  $i == j$  or  $(i, j) \in G.E$ 
6               $t_{ij}^{(0)} = 1$ 
7          else  $t_{ij}^{(0)} = 0$ 
8  for  $k = 1$  to  $n$ 
9      let  $T^{(k)} = (t_{ij}^{(k)})$  be a new  $n \times n$  matrix
10     for  $i = 1$  to  $n$ 
11         for  $j = 1$  to  $n$ 
12              $t_{ij}^{(k)} = t_{ij}^{(k-1)} \vee (t_{ik}^{(k-1)} \wedge t_{kj}^{(k-1)})$ 
13 return  $T^{(n)}$ 
```


Johnson's Algorithm

Johnson's Algorithm computes all-pair shortest paths for all the vertices.

- Suitable for **sparse graphs**
- Asymptotically better than the **Floyd-Warshall** algorithm
- Either returns a matrix of shortest paths for all pairs of vertices or reports that the graph contains a **negative-weight cycle**
- Uses both **Dijkstra's algorithm** and the **Bellman-Ford** algorithm **as subroutines**
- Uses the **reweighting technique**

Johnson's Algorithm: Reweighting

If all weights w in a graph $G = (V, E, w)$ are **non-negative**, we can find shortest paths between all pairs of vertices by running **Dijkstra's algorithm** on **each vertex** in the graph.

- The running time of this part is $O((V^2 + VE) \log V)$ with a **min-priority Q**.

If G has **negative-weight edges** but **no negative-weight cycles**, we simply compute the new set of edge weights so that we can apply **Dijkstra's algorithm**.

- The new set of edge weights \hat{w} must satisfy the following **two properties**:
 - **(P1)** For all pairs of vertices $u, v \in V$, a path p is a shortest path from u to v using weight function w if and only if p is also a shortest path from u to v using weight function \hat{w}
 - **(P2)** For all edges (u, v) , the new weight $\hat{w}(u, v)$ is **non-negative**

Johnson's Algorithm: Reweighting

Lemma: (P1: Reweighting does not change shortest paths)

Given a **weighted, directed** graph $G = (V, E, w)$ with a weight function $w: E \rightarrow \mathbb{R}$, let $h: V \rightarrow \mathbb{R}$ be a function mapping vertices to real numbers, for each edge $(u, v) \in E$, we define

$$\hat{w}(u, v) = w(u, v) + h(u) - h(v). \quad (\text{Eq.7})$$

Let $p = \langle v_0, v_1, v_2, \dots, v_k \rangle$ be any path from v_0 to v_k .

Then, p is a shortest path from v_0 to v_k if and only if p is also a shortest path with weight function \hat{w} .

(Claim I) That is, $w(p) = \delta(v_0, v_k)$ if and only if $\hat{w}(p) = \hat{\delta}(v_0, v_k)$. [$\hat{\delta}$ denotes shortest-path weights derived from weight function \hat{w} .]

(Claim II) Further more, G has a negative-weight cycle using w if and only if G has a negative-weight cycle using \hat{w} .

Johnson's Algorithm: Reweighting

Proof:

(Claim I)

We will start by showing that $\hat{w}(p) = w(p) + h(v_0) - h(v_k)$.

Since

$$\begin{aligned}\hat{w}(p) &= \sum_{i=1}^k \hat{w}(v_{i-1}, v_i) \\ &= \sum_{i=1}^k (w(v_{i-1}, v_i) + h(v_{i-1}) - h(v_i)) \\ &= \sum_{i=1}^k (w(v_{i-1}, v_i) + h(v_0) - h(v_k)) \\ &= w(p) + h(v_0) - h(v_k)\end{aligned}$$

Since $h(v_0)$ and $h(v_k)$ are **path-independent**, if one path from v_0 to v_k is shorter than another using weight function w , then it is also shorter using weight function \hat{w} .

Thus, we have that $w(p) = \delta(v_0, v_k)$ if and only if $\hat{w}(p) = \hat{\delta}(v_0, v_k)$. ■

Johnson's Algorithm: Reweighting

Proof:

(Claim II)

We show that G has a negative-weight cycle using w if and only if G has a negative-weight cycle using \hat{w} .

Consider any cycle $c = \langle v_0, v_1, v_2, \dots, v_k \rangle$, where $v_0 = v_k$.

$$\hat{w}(c) = w(c) + h(v_0) - h(v_k) \quad [\text{Eq. 7}]$$

$$\hat{w}(c) = w(c) \quad [v_0 = v_k \rightarrow h(v_0) = h(v_k)]$$

Therefore, G has a negative-weight cycle using w if and only if G has a negative-weight cycle using \hat{w} . ■

Johnson's Algorithm: Reweighting

Lemma: (P2: Reweighting ensures non-negativity)

Given a **weighted, directed** graph $G = (V, E, w)$ with a weight function $w: E \rightarrow \mathbb{R}$, let $h: V \rightarrow \mathbb{R}$ be a function mapping vertices to real numbers. we define

$$\hat{w}(u, v) = w(u, v) + h(u) - h(v).$$

For all edges (u, v) , the new weight $\hat{w}(u, v)$ is **non-negative**.

Proof:

We construct a new graph $G' = (V', E')$, where $V' = V \cup \{s\}$ for some new vertex $s \notin V$ and $E' = E \cup \{(s, v): v \in V\}$.

We extend the weight function w so that $w(s, v) = 0$ for all $v \in V$.

Observation: G' has **no negative-weight cycles** if and only if G has **no negative-weight cycles**.

Johnson's Algorithm: Reweighting

Proof: (Continued)

Suppose that G and G' have **no negative-weight cycles**.

Define $h(v) = \delta(s, v)$ for all $v \in V'$.

$$h(v) \leq h(u) + w(u, v) \text{ for all } (u, v) \in E' \quad [\text{Triangle Inequality}] \text{ (See Lecture 10)}$$


$$h(v) - h(u) \leq w(u, v)$$

$$0 \leq w(u, v) + h(u) - h(v)$$


$$0 \leq \hat{w}(u, v) \quad [\text{Eq. 7}]$$

For all edges (u, v) , the new weight $\hat{w}(u, v)$ is **non-negative**. ■


Johnson's Algorithm: Example




(a)




(b)




(c)




(d)



(e)



(f)



(g)

Johnson's Algorithm

JOHNSON(G, w)

```
1  compute  $G'$ , where  $G'.V = G.V \cup \{s\}$ ,  
    $G'.E = G.E \cup \{(s, v) : v \in G.V\}$ , and  
    $w(s, v) = 0$  for all  $v \in G.V$   
2  if BELLMAN-FORD( $G', w, s$ ) == FALSE  
3    print "the input graph contains a negative-weight cycle"  
4  else for each vertex  $v \in G'.V$   
5    set  $h(v)$  to the value of  $\delta(s, v)$   
   computed by the Bellman-Ford algorithm  
6  for each edge  $(u, v) \in G'.E$   
7     $\hat{w}(u, v) = w(u, v) + h(u) - h(v)$   
8  let  $D = (d_{uv})$  be a new  $n \times n$  matrix  
9  for each vertex  $u \in G.V$   
10   run DIJKSTRA( $G, \hat{w}, u$ ) to compute  $\hat{\delta}(u, v)$  for all  $v \in G.V$   
11   for each vertex  $v \in G.V$   
12      $d_{uv} = \hat{\delta}(u, v) + h(v) - h(u)$   
13  return  $D$ 
```

Johnson's Algorithm: Analysis

The running time of Johnson's algorithm is determined by the for loop of **lines 9-12**.

Running Dijkstra's algorithm V times costs $O((V^2 + VE) \log V)$.

If G is sparse,

then $E = O(V)$ so the running time becomes $O((V^2 + V \cdot V) \log V) = O(V^2 \log V)$, which is asymptotically faster than the **Floyd-Warshall** algorithm, which takes $\Theta(V^3)$ time.

Summary

In this lecture, we have covered the topic of all-pair shortest path problems:

- All-Pair Shortest Path Problems
 - Naive Dynamic Programming View
 - Matrix Multiplication View
 - Floyd-Warshall Algorithm
 - Transitive Closure of Directed Graphs
 - Johnson's Algorithm